

Rebuilding Slack Search with Armeria

November 21, 2019





Renaud Bourassa

Staff Software Engineer
Search & Discovery

Slack

Our mission is to make people's working lives simpler, more pleasant, and more productive.


The screenshot displays the Slack application interface. On the left is a dark purple sidebar for the workspace 'Acme Inc.' (owner: Alex Saunders). It contains a 'Jump to...' search bar, 'All Unreads', 'Starred' items (announcements, design-team with 1 unread), and a list of channels: #social-media (selected), #helpdesk, and #Sara Parker. Below these are 'Channels' (accounting, design-crit, help-design, media-and-pr, triage-issues, design-team-sf) and 'Direct Messages' (slackbot, Zoe Maxwell, Leland..., Florence Garret, Liza Zhang). The main area shows the '#social-media' channel with a header indicating 21 members and a description 'Track and coordinate social med...'. The chat history includes messages from Sara Parker (12:53 PM) praising @zoe, Zoe Maxwell (12:55 PM) expressing pleasure, an event announcement from Acme Team BOT (12:56 PM) for a 'Team Status Meeting' from 1:00 PM to 1:30 PM, a quick note from Harry Boone (12:58 PM) about @Liza joining a sync, and meeting notes from Jeremy Stevens (1:42 PM). A '1/9 Meeting Notes' document is shown as a post. The right sidebar, titled 'About #social-media', lists channel details, highlights, 1 pinned item, 21 members, shared files, and notification preferences. The bottom of the interface features a message input field with a placeholder 'Message #social-media' and icons for attachments, mentions, and emojis.

Stats

12M+ daily active users (half outside of North America)

5B+ actions/week

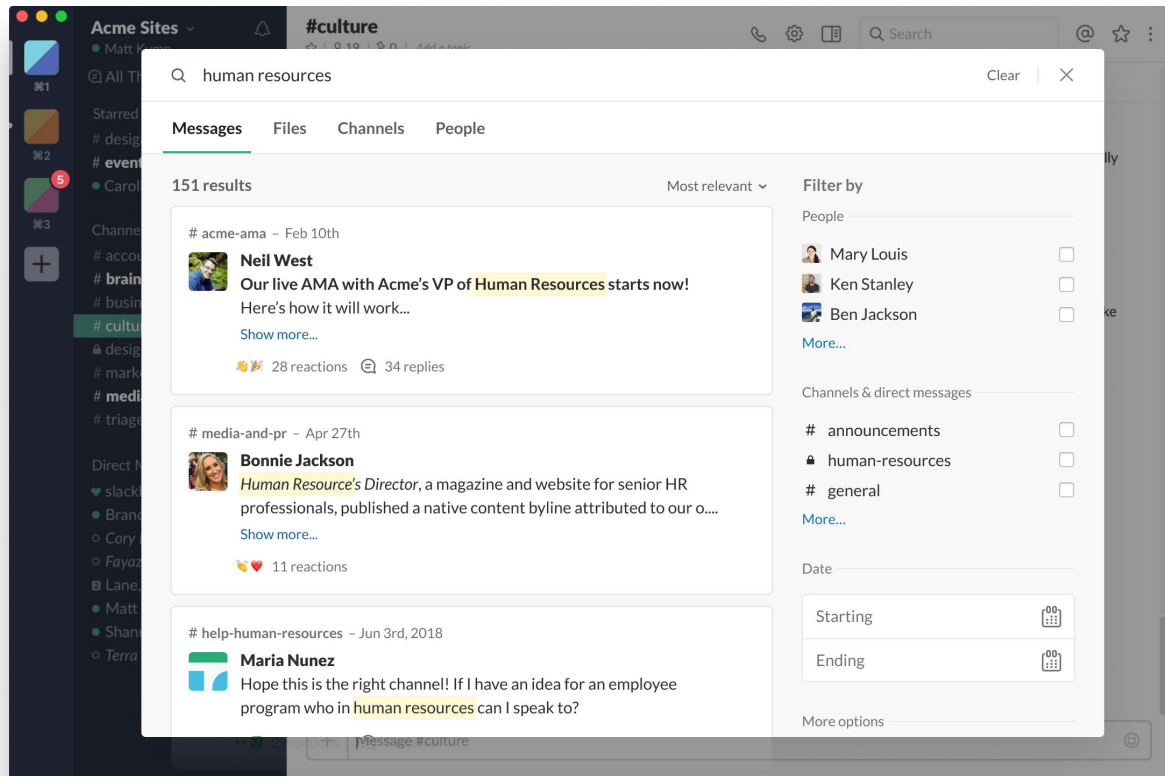
Daily Active Users



In September 2019, we exceeded 12 million daily active users

Search & Discovery

We provide tools to help people find, discover, organize and focus on the information they need to get work done in Slack.



A Brief History of Search

August 2013

Slack open beta
launches with Search
as a core feature.

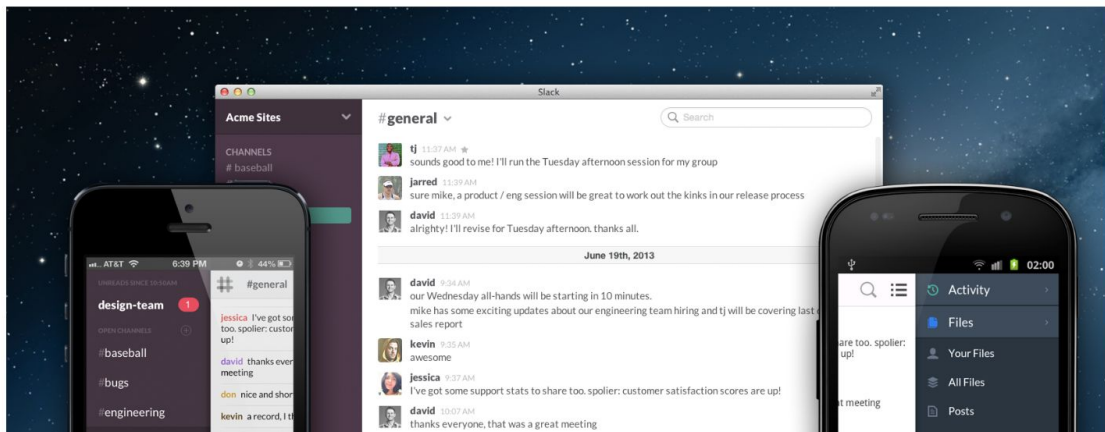
Slack, The Newest Enterprise Social Network, Is The Latest Effort From Flickr Co-Founder Stewart Butterfield



Ingrid Lunden @ingridlunden / 11:18 am EDT • August 14, 2013



Comment



November 2012

Search was always seen as core to the Slack experience.



Stewart Butterfield ✓

@stewart



Yes. November 14th, 2012 (previous codename was "linefeed"):

Re [twitter.com/minney_cat/sta...](#)

9:48 AM <stewart> I was up very late last night thinking about things

9:49 AM <stewart> and it is no exaggeration to say "WOW!"

9:49 AM <stewart> Also, I have a better code name

9:49 AM <stewart> "Slack" and/or "Slack App"

9:49 AM <stewart> **Searchable Log of All Conversation & Knowledge**

9:49 AM <stewart> but, "slack" is also just nice to say

9:51 AM <eric> I like it, but... it has kind of negative connotations, too.

9:51 AM <eric> Our users would be Slackers :)

9:52 AM <stewart> ta da!

9:52 AM <stewart> just a codename

9:52 AM <eric> ah, cool

9:53 AM <stewart> but wait till I have Kuke draw you the "slack" (8 armed, fuzzy-headed, ...)

Minn @minney_cat

@stewart is this backronym for @SlackHQ true? Searchable Log of All Conversation and Knowledge 🤔

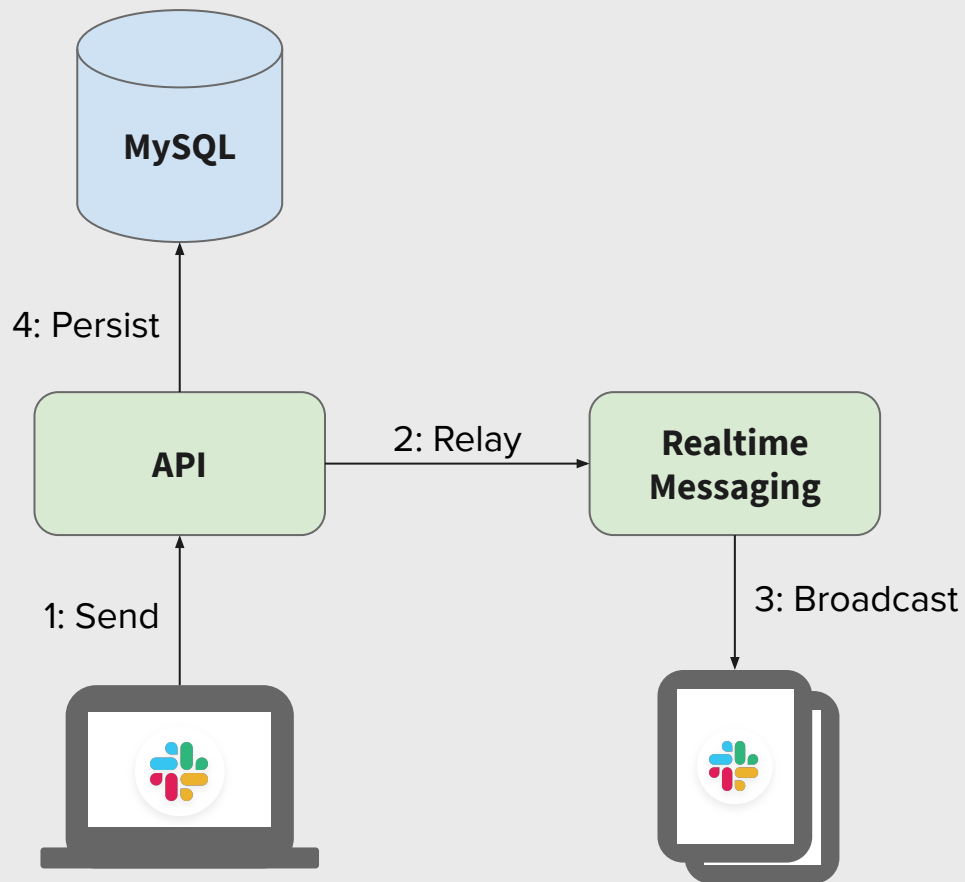
♥ 1,094 7:07 PM - Sep 27, 2016



💬 395 people are talking about this

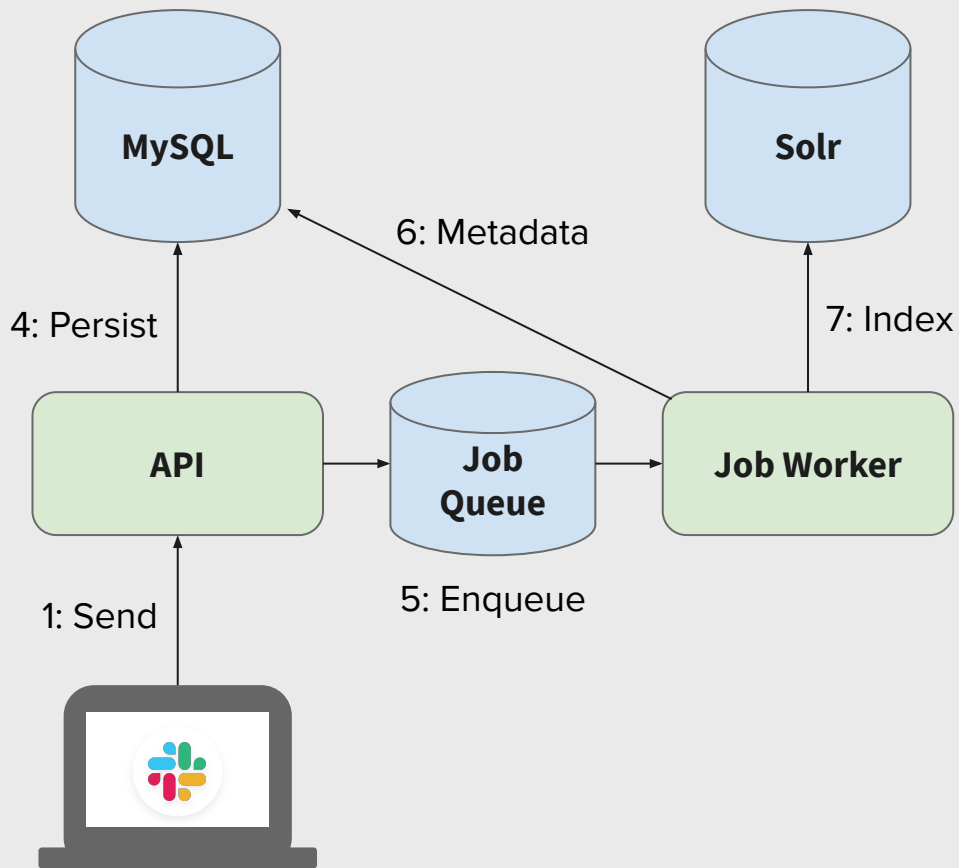


Slack Architecture



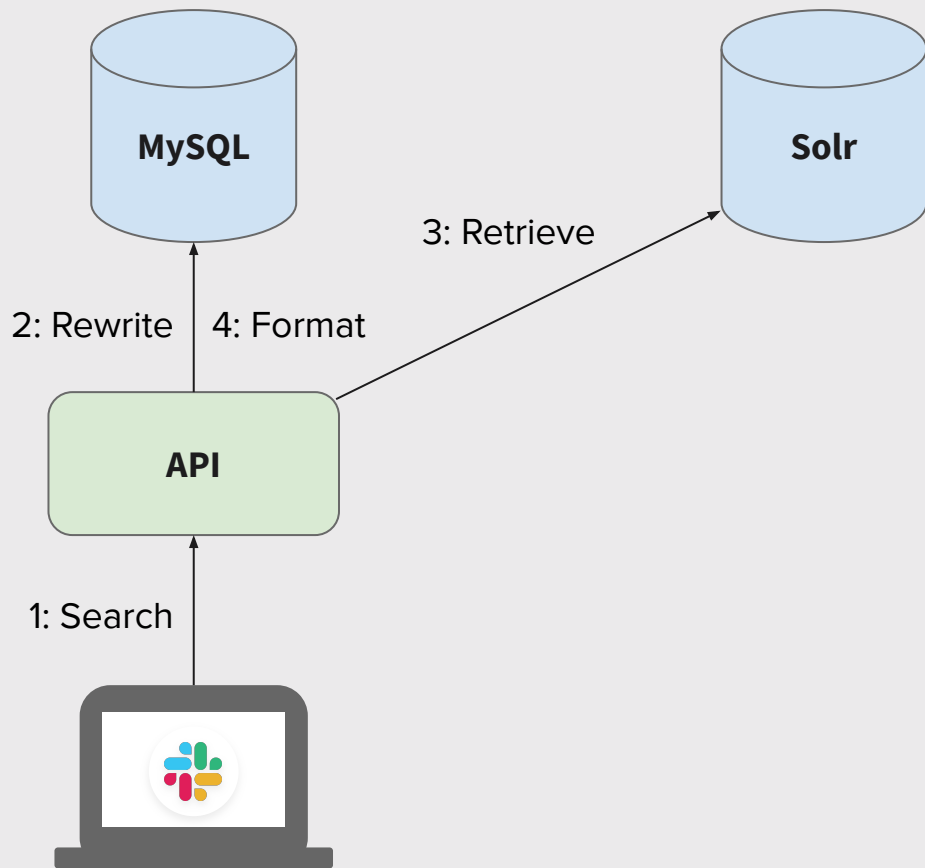
Search Indexing

Write Path



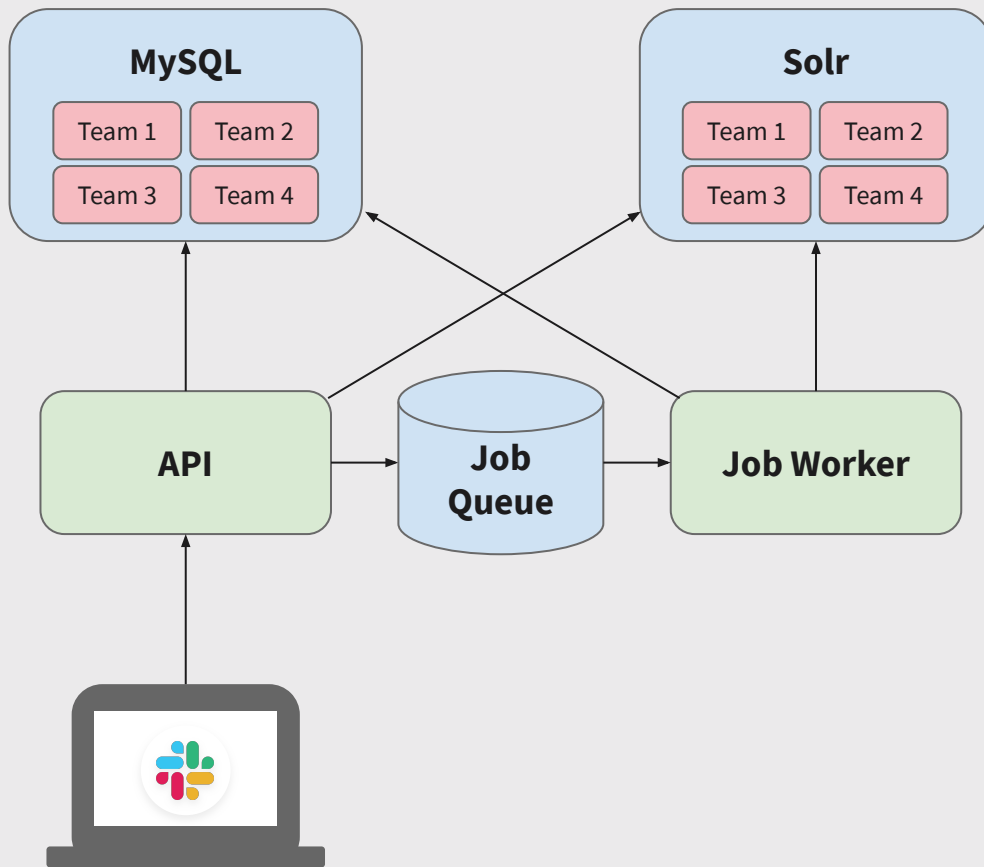
Search Querying

Read Path



Scaling

Team Sharding




**A Few
Years Later**

Indexing

**Search
performance is
getting worse.**

Index size per shard (November 2016)




Querying

Search quality is getting worse.

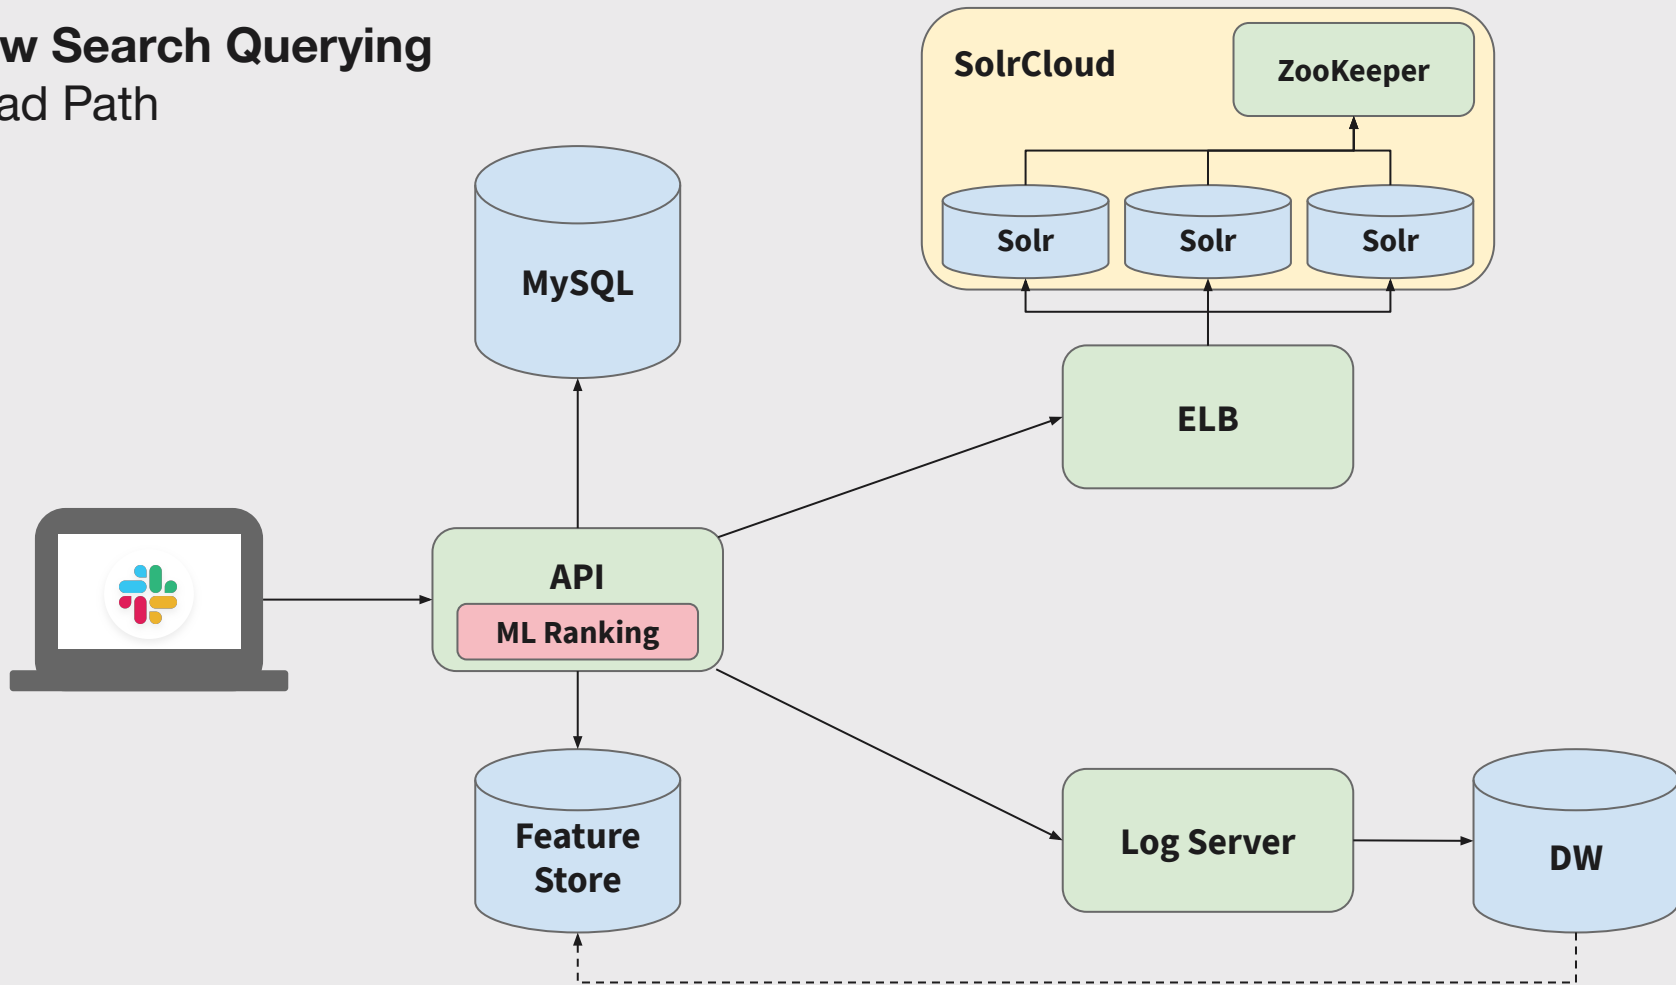
Score = TF-IDF

```
log(numDocs / (docFreq + 1)) * sqrt(tf) * (1/sqrt(length))
```

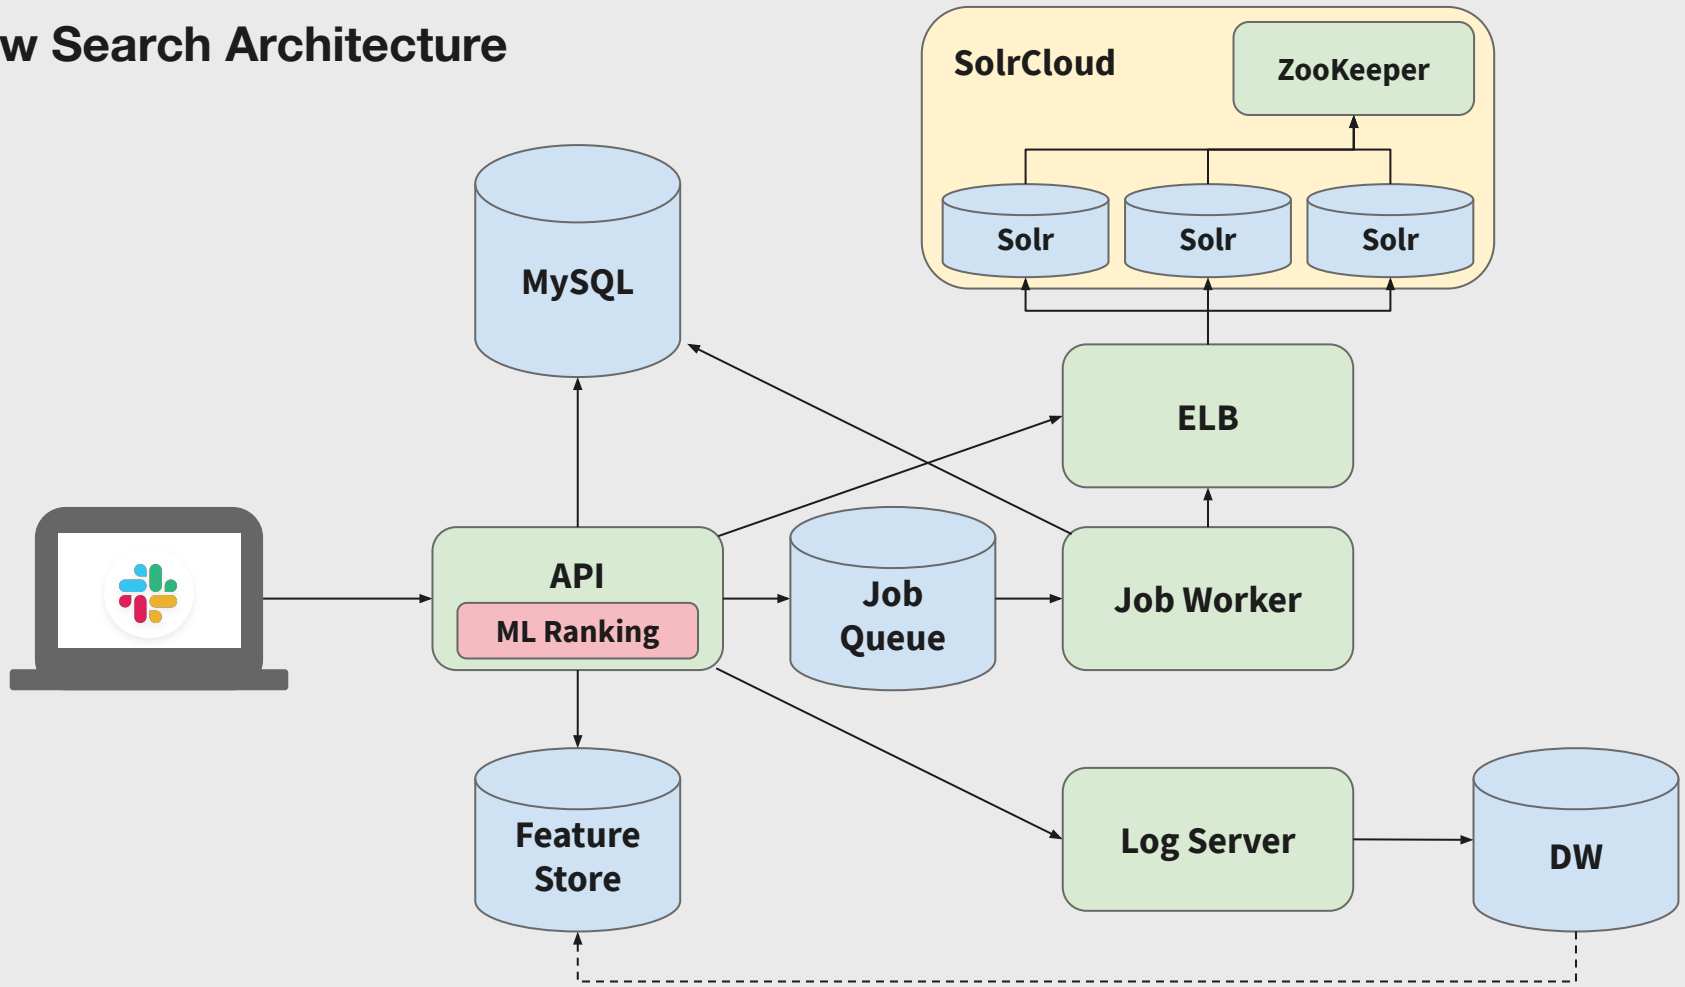
New Search Indexing Read Path



New Search Querying Read Path

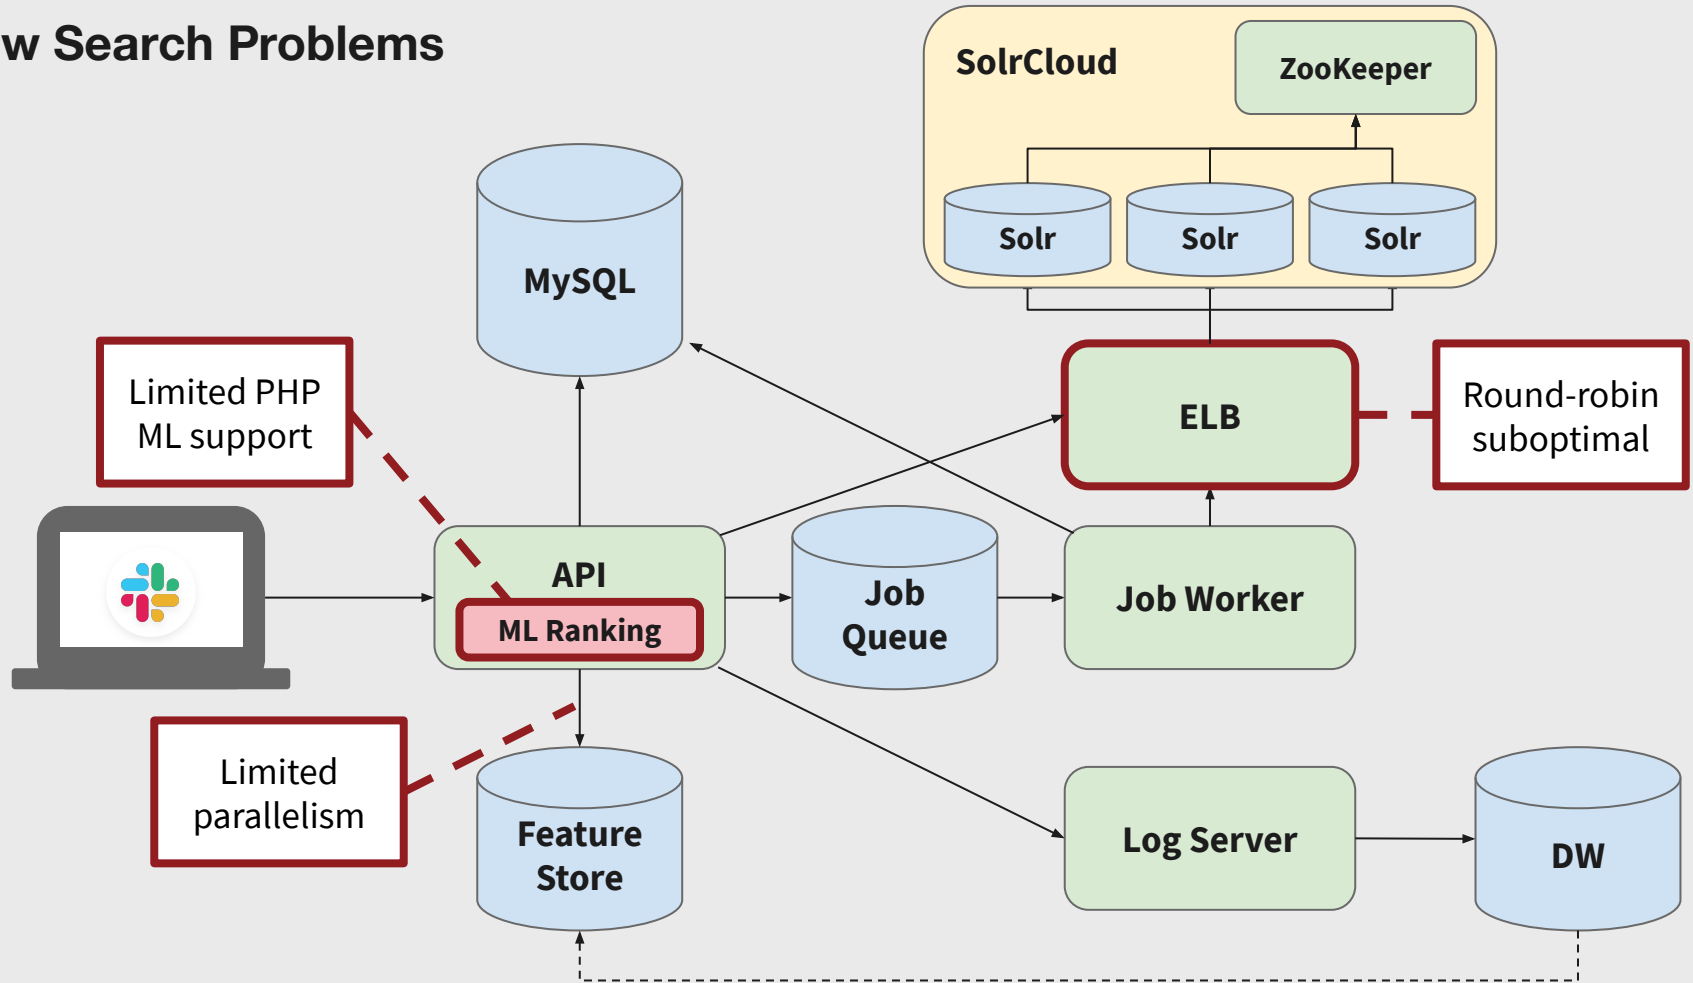


New Search Architecture

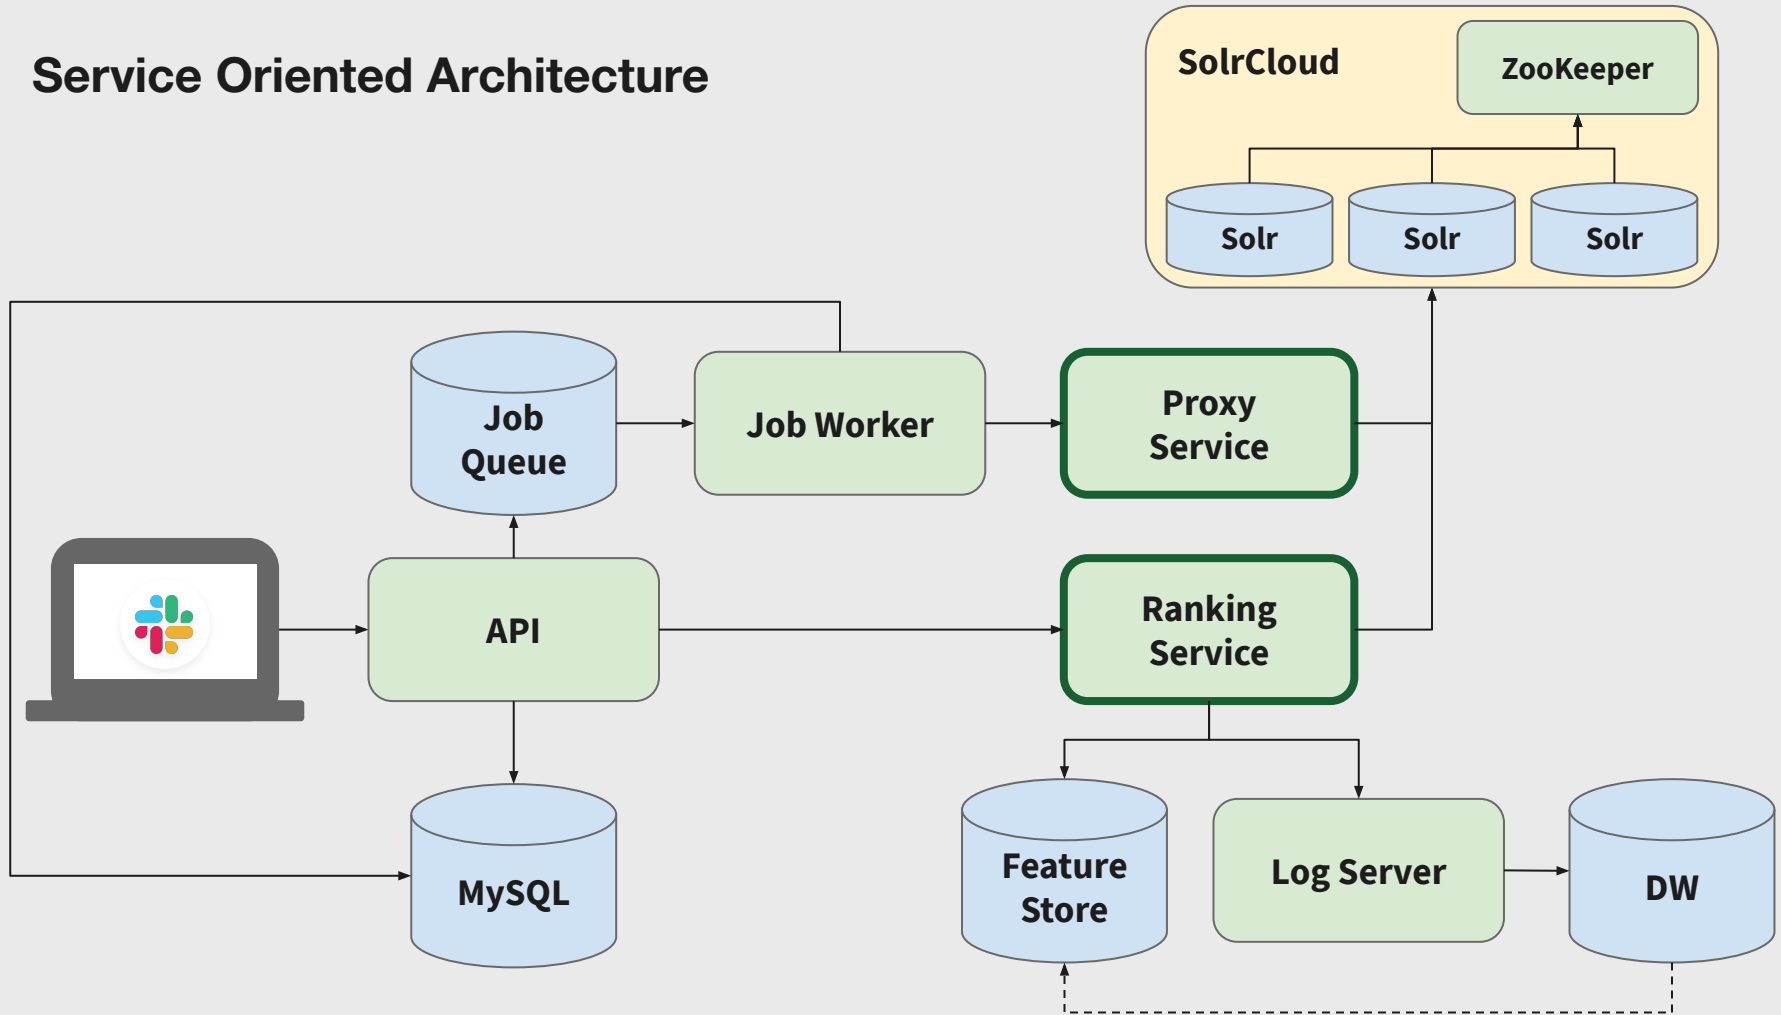


**A Few
Months Later**

New Search Problems



Service Oriented Architecture



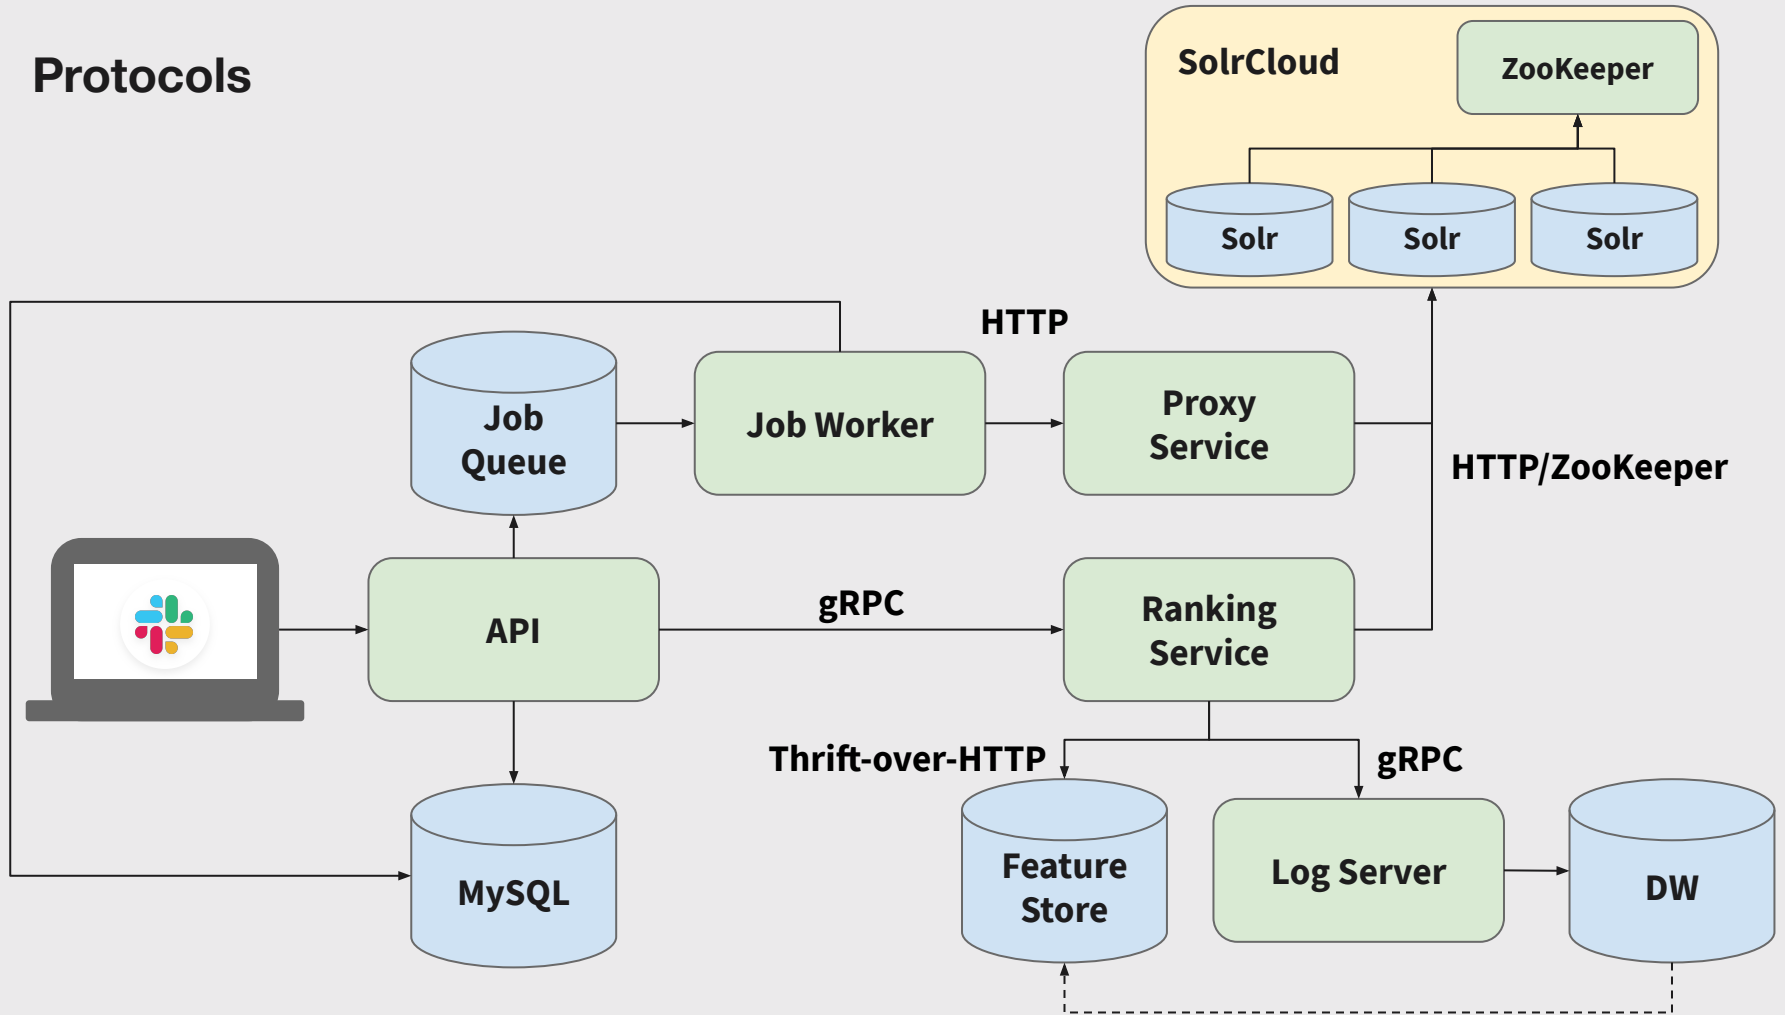
Requirements

- **Written in Java**
 - Solr, SolrJ, XGBoost4J

Requirements

- **Written in Java**
 - Solr, SolrJ, XGBoost4J
- **Supports multiple protocols**
 - HTTP, gRPC, Thrift-over-HTTP

Protocols



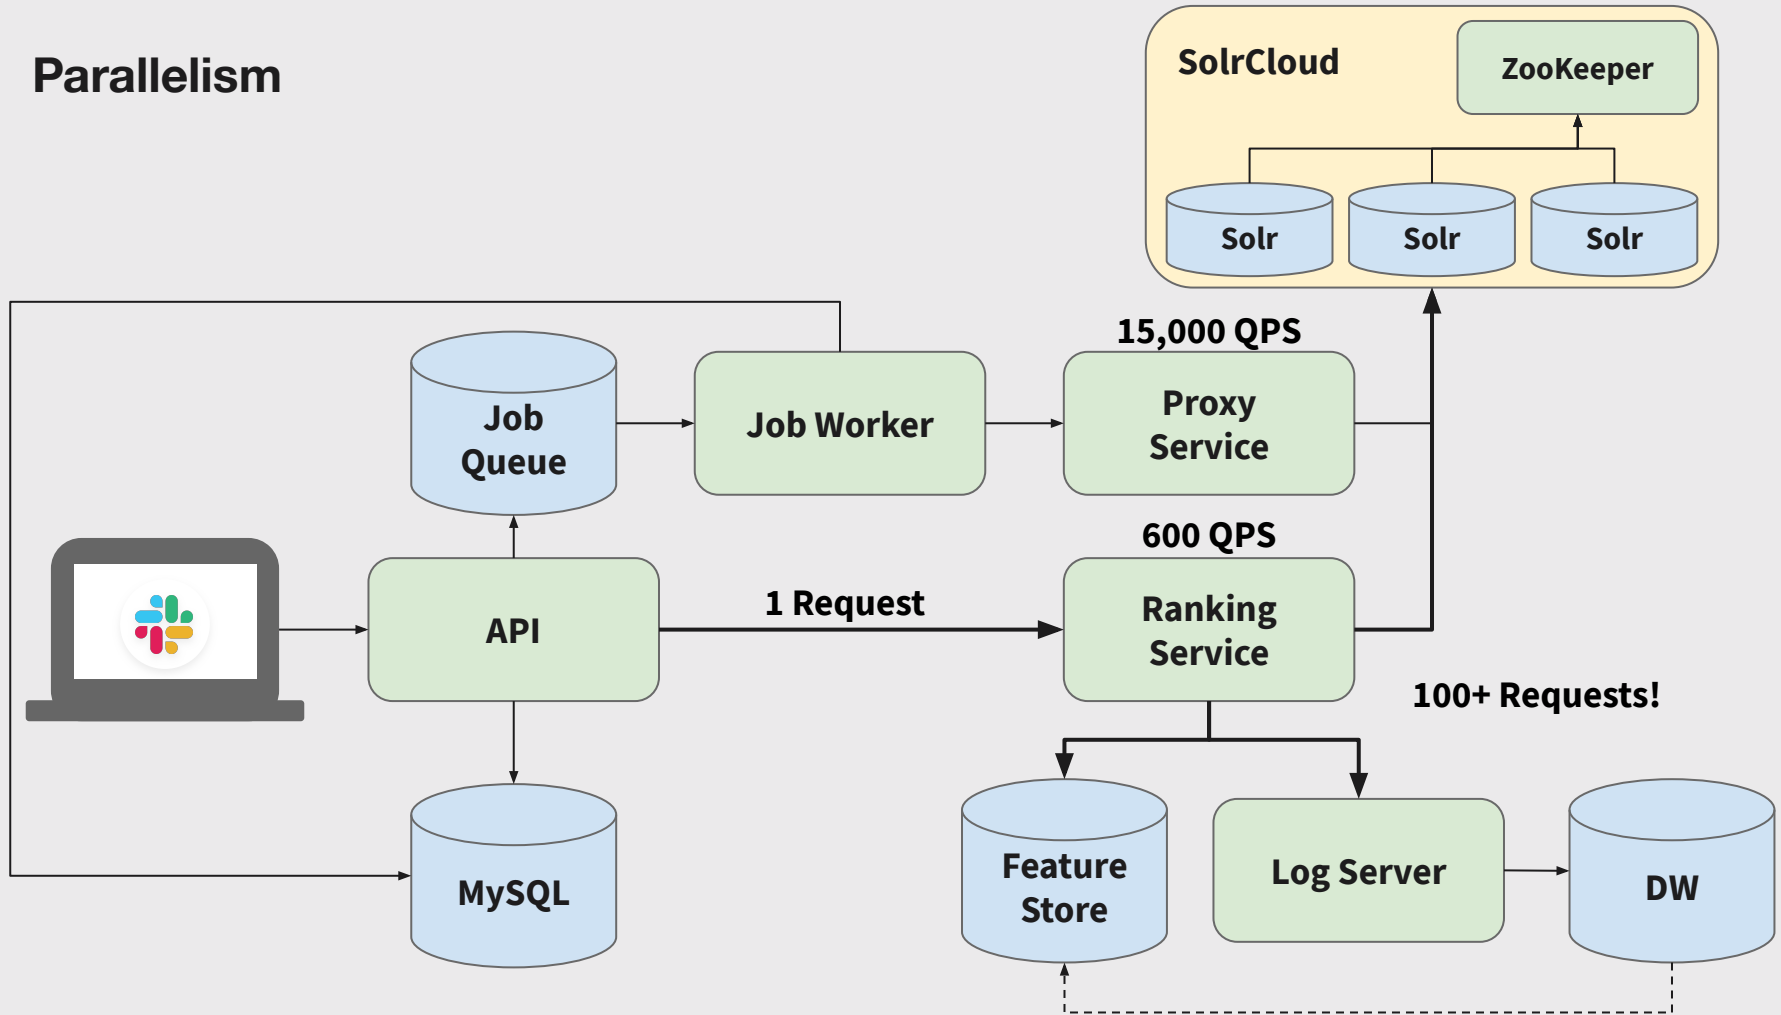
Requirements

- **Written in Java**
 - Solr, SolrJ, XGBoost4J
- **Supports multiple protocols**
 - HTTP, gRPC, Thrift-over-HTTP
- **Integrates with existing tools**
 - Consul, Prometheus, Rsyslog, Honeycomb

Requirements

- **Written in Java**
 - Solr, SolrJ, XGBoost4J
- **Supports multiple protocols**
 - HTTP, gRPC, Thrift-over-HTTP
- **Integrates with existing tools**
 - Consul, Prometheus, Rsyslog, Honeycomb
- **Good parallel performance**

Parallelism



Requirements

- **Written in Java**
 - Solr, SolrJ, XGBoost4J
- **Supports multiple protocols**
 - HTTP, gRPC, Thrift-over-HTTP
- **Integrates with existing tools**
 - Consul, Prometheus, Rsyslog, Honeycomb
- **Good parallel performance**



Armeria

Armeria

Armeria is an open-source asynchronous HTTP/2 RPC/REST client/server library built on top of Java 8, Netty, Thrift and gRPC. Its primary goal is to help engineers build high-performance asynchronous microservices that use HTTP/2 as a session layer protocol.

It is open-sourced and licensed under Apache License 2.0 by LINE Corporation, who uses it in production.

Requirements



Written in Java

- Solr, SolrJ, XGBoost4J


• Supports multiple protocols

- HTTP, gRPC, Thrift-over-HTTP

• Integrates with existing tools

- Consul, Prometheus, Rsyslog, Honeycomb

• Good parallel performance



Armeria

<http://line.github.io/armeria/>

Server (HTTP)

```
public class ProxyServlet extends HttpServlet {  
    protected void doPost(HttpServletRequest req, HttpServletResponse resp)  
        throws ServletException, IOException {  
        ...  
    }  
}
```

```
ServletContextHandler handler = new ServletContextHandler(NO_SESSIONS);  
handler.setClassLoader(SolrCloudProxyServer.class.getClassLoader());  
handler.addServlet(ProxyServlet.class, "/*");  
Server arm = new ServerBuilder()  
    .serviceUnder("/solr", new JettyServiceBuilder().handler(handler).build())  
    .build();  
arm.start();
```


Server (gRPC)

```
service SolrRankService {  
    rpc Rank(SolrRankRequest) returns (SolrRankResponse) {}  
}
```

```
public class SolrRankService extends SolrRankServiceGrpc.SolrRankServiceImplBase {  
    public void rank(SolrRankRequest req, StreamObserver<SolrRankResponse> responseObserver) {  
        ...  
    }  
}
```

```
SolrRankService solrRankService = new SolrRankService(...);  
Server arm = new ServerBuilder()  
    .service(new GrpcServiceBuilder().addService(solrRankService).build())  
    .build();  
arm.start();
```

Client (Thrift)

```
service HFileService {  
    SingleHFileKeyResponse getValuesSingle(1: SingleHFileKeyRequest req)  
    ...  
}
```

```
HFileService.Iface client =  
    new ClientBuilder("tbinary+http://...")  
        .writeTimeout(DEFAULT_CONNECTION_TIMEOUT)  
        .decorator(  
            RetryingHttpClient.newDecorator(  
                RetryStrategy.onServerErrorStatus(DEFAULT_BACKOFF),  
                DEFAULT_MAX_TOTAL_ATTEMPTS,  
                DEFAULT_RESPONSE_TIMEOUT.toMillis()));  
        .build(HFileService.Iface.class);
```

Requirements



Written in Java

- Solr, SolrJ, XGBoost4J




Supports multiple protocols

- HTTP, gRPC, Thrift-over-HTTP

● Integrates with existing tools

- Consul, Prometheus, Rsyslog, Honeycomb

● Good parallel performance




Armeria

<http://line.github.io/armeria/>

Decorators

```
new ServerBuilder()  
  .decorator((delegate, ctx, req) -> {  
    // Do something with server request.  
    return delegate.execute(ctx, req);  
  })  
  .build();
```

```
new ClientBuilder(...)  
  .decorator((delegate, ctx, req) -> {  
    // Do something with client request.  
    return delegate.execute(ctx, req);  
  })  
  .build(...);
```



Logging

```
Server arm = new ServerBuilder()
    ...
    .decorator(
        new LoggingServiceBuilder()
            .requestLogLevel(LogLevel.INFO);
            .successfulResponseLogLevel(LogLevel.INFO)
            .failureResponseLogLevel(LogLevel.ERROR)
            .requestHeadersSanitizer(
                headers -> {
                    if (headers.contains(HttpHeaderNames.AUTHORIZATION)) {
                        return headers
                            .toBuilder()
                            .removeAndThen(HttpHeaderNames.AUTHORIZATION)
                            .add(HttpHeaderNames.AUTHORIZATION, "present_but_scrubbed")
                            .build();
                    }
                    return headers;
                })
            .newDecorator())
    ...
```

Logging

🔍 🔍 📄 * I 1030 21:11:52.506899 7082 com.linecorp.armeria.server.HttpServerHandler:415 \t\t[id: 0x3dbccc75, L:/[REDACTED] - R:/[REDACTED]][h1c://[REDACTED]]
[REDACTED]/solr/mc-computed_v201910180000_0.1/update#POST] Request: {startTime=2019-10-30T21:11:52.506Z(1572469912506000), length=591B, duration=31260ns, scheme=none+h1c, headers=[:method=POST, :scheme=http, :authority=[REDACTED], :path=/solr/mc-computed_v201910180000_0.1/update?version=2.2&indent=on&tw=&wt=phps, accept-encoding=gzip,deflate, user-agent=Slack SolrClient 1.0.1, accept=/*/*, accept-charset=utf-8, content-type=text/xml; charset=UTF-8, content-length=591, authorization=present_but_scrubbed]}

Metrics Prometheus

```
Server arm = new ServerBuilder()  
    ...  
    .service(  
        "/metrics",  
        new PrometheusExpositionService(  
            config.getMeterRegistry().getPrometheusRegistry()  
        )  
    )  
    ...  
    .decorator(  
        MetricCollectingService.newDecorator(  
            MeterIdPrefixFunction.ofDefault(config.getMeterPrefix())  
                .withTags(config.getMeterTags()))  
    )  
    ...
```

Metrics Prometheus

```
# HELP solrjproxy_requests_total
# TYPE solrjproxy_requests_total counter
solrjproxy_requests_total{hostnamePattern="*",httpStatus="500",method="POST",result="success",route="prefix:/solr/",version="948"}, 0.0
solrjproxy_requests_total{hostnamePattern="*",httpStatus="200",method="GET",result="failure",route="prefix:/health/",version="948"}, 0.0
solrjproxy_requests_total{hostnamePattern="*",httpStatus="0",method="POST",result="failure",route="prefix:/solr/",version="948"}, 123.0
solrjproxy_requests_total{hostnamePattern="*",httpStatus="401",method="POST",result="failure",route="prefix:/solr/",version="948"}, 17.0
solrjproxy_requests_total{hostnamePattern="*",httpStatus="200",method="GET",result="success",route="exact:/metrics",version="948"}, 36419.0
solrjproxy_requests_total{hostnamePattern="*",httpStatus="200",method="GET",result="failure",route="exact:/metrics",version="948"}, 0.0
solrjproxy_requests_total{hostnamePattern="*",httpStatus="500",method="POST",result="failure",route="prefix:/solr/",version="948"}, 242.0
solrjproxy_requests_total{hostnamePattern="*",httpStatus="200",method="POST",result="success",route="prefix:/solr/",version="948"}, 3.2823618E7
solrjproxy_requests_total{hostnamePattern="*",httpStatus="401",method="GET",result="failure",route="prefix:/solr/",version="948"}, 16.0
solrjproxy_requests_total{hostnamePattern="*",httpStatus="200",method="GET",result="success",route="prefix:/solr/",version="948"}, 769593.0
solrjproxy_requests_total{hostnamePattern="*",httpStatus="503",method="GET",result="failure",route="prefix:/solr/",version="948"}, 1.0
```

```
# HELP solrjproxy_total_duration_seconds
# TYPE solrjproxy_total_duration_seconds summary
solrjproxy_total_duration_seconds{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948",quantile="0.0"}, 0.006291456
solrjproxy_total_duration_seconds{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948",quantile="0.5"}, 0.00917504
solrjproxy_total_duration_seconds{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948",quantile="0.75"}, 0.009699328
solrjproxy_total_duration_seconds{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948",quantile="0.9"}, 0.010223616
solrjproxy_total_duration_seconds{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948",quantile="0.95"}, 0.010223616
solrjproxy_total_duration_seconds{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948",quantile="0.98"}, 0.020709376
solrjproxy_total_duration_seconds{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948",quantile="0.99"}, 0.020709376
solrjproxy_total_duration_seconds{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948",quantile="0.999"}, 0.020709376
solrjproxy_total_duration_seconds{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948",quantile="1.0"}, 0.020709376
solrjproxy_total_duration_seconds_count{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948"}, 32976.0
solrjproxy_total_duration_seconds_sum{hostnamePattern="*",httpStatus="409",method="POST",route="prefix:/solr/",version="948"}, 333.385478434
```

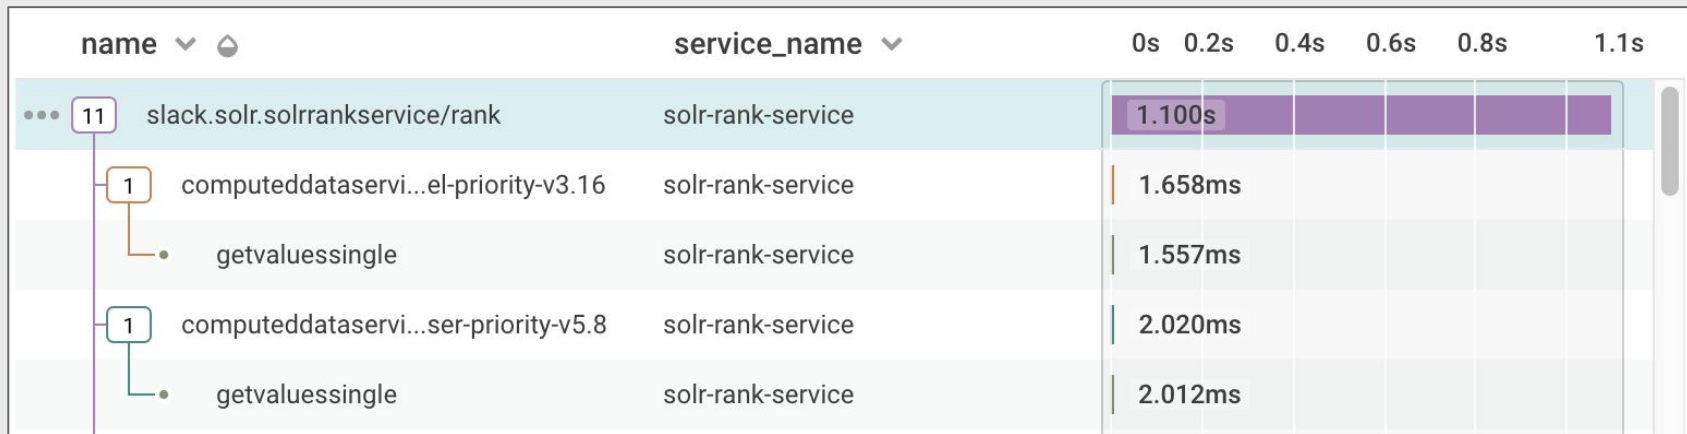

Metrics Prometheus

```
HFileService.Iface client =  
    new ClientBuilder("tbinary+http://...")  
        ...  
        .decorator(  
            MetricCollectingClient.newDecorator(MeterIdPrefixFunction.ofDefault("cds"))  
        )  
        ...
```



Tracing honeycomb

```
Server arm = new ServerBuilder()  
...  
.decorator(BraveService.newDecorator(config.getTracing()))  
...
```



Tracing honeycomb

```
public class MetricCollectingSolrClient extends SolrClient {

    private MetricCollectingSolrClient(SolrClient solrClient, Tracing tracing) {
        this.solrClient = solrClient;
        this.tracing = tracing;
    }

    public NamedList<Object> request(SolrRequest request, String collection)
        throws SolrServerException, IOException {
        Span span = tracing.tracer().newChild(RequestContextCurrentTraceContext.ofDefault().get());
        span.name(getName(collection));
        try (SpanInScope ignore = tracing.tracer().withSpanInScope(span)) {
            try {
                return solrClient.request(request, collection);
            } finally {
                span.finish();
            }
        }
    }
    ...
}
```

Service Discovery Consul

```
Server arm = new ServerBuilder()
...
.service("/health", new ManagedHttpHealthCheckService())
...
.serverListener(
    ConsulUpdatingListener.builder(config.getServerName(), config.getPort())
        .serviceId(config.getServerName() + ":" + config.getPort())
        .address(config.getServerAddress())
        .check(String.format("http://localhost:%d/health/", config.getPort()), 10, 5, 15)
        .tag(String.format("version:%s", config.getVersion()))
        .build())
...

```

Service Discovery Consul

HashiCorp

```
public class ConsulUpdatingListener extends ServerListenerAdapter {

    public void serverStarted(Server server) throws Exception {
        ImmutableRegistration.Builder registrationBuilder =
            ImmutableRegistration.builder()
                .address(address)
                .port(port)
                .checks(checks)
                .name(service)
                .id(serviceId)
                .addTags(tags)
        }
        agentClient.register(registrationBuilder.build());
    }

    public void serverStopping(Server server) throws Exception {
        agentClient.deregister(serviceId);
        Thread.sleep(15000);
    }


    ...
}
```

Service Discovery HashiCorp Consul

Why sleep on stop? **To drain requests!**

```
Runtime.getRuntime()
    .addShutdownHook(
        new Thread() {
            @Override
            public void start() {
                log.info("Shutdown hook called");
                arm.stop();
            }
        });
```

Service Discovery HashiCorp Consul



<https://line.github.io/armeria/client-service-discovery.html>

```
public static void registerEndpoint(  
    String service, String partition, int refreshEvery, TimeUnit unit) {  
    EndpointGroupRegistry.register(  
        service,  
        new ConsulEndpointGroup(service, partition, refreshEvery, unit),  
        EndpointSelectionStrategy.ROUND_ROBIN);  
}
```

```
new ClientBuilder(String.format("tbinary+http://group:%s/", service))
```

Service Discovery Consul

HashiCorp

```
public class ConsulEndpointGroup extends DynamicEndpointGroup {

    public ConsulEndpointGroup(String service, String dataCenter, int refreshEvery, TimeUnit unit) {
        ...
        this.exec = Executors.newSingleThreadScheduledExecutor();
        exec.scheduleAtFixedRate(this::refresh, 0, refreshEvery, unit);
    }

    void refresh() {
        List<Endpoint> newData =
            healthClient
                .getHealthyServiceInstances(service)
                .getResponse();
                .stream()
                .map(n -> Endpoint.of(n.getService().getAddress(), n.getService().getPort()))
                .collect(Collectors.toList());
        if (!endpoints().equals(newData)) {
            setEndpoints(newData);
        }
    }
    ...
}
```


Requirements



Written in Java

- Solr, SolrJ, XGBoost4J



Supports multiple protocols


- HTTP, gRPC, Thrift-over-HTTP



Integrates with existing tools

- Consul, Prometheus, Rsyslog, Honeycomb

- **Good parallel performance**



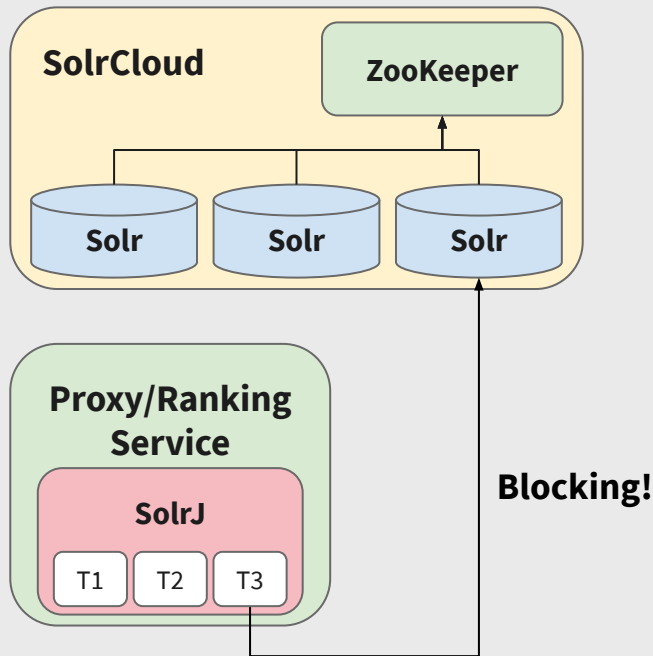
Armeria

<http://line.github.io/armeria/>

Performance


name ▼		service_name ▼	0s	0.2s	0.4s	0.6s	0.7884s
11	slack.solr.solrankservice/rank	solr-rank-service	788.4ms				
3	computeddataservi...el-priority-v3.16	solr-rank-service	4.050ms				
3	computeddataservi...ser-priority-v5.8	solr-rank-service	1.435ms				
3	computeddataservi...ser-affinity-v2.8	solr-rank-service	3.400ms				
51	metriccollectingsolrclient:chat_8.64	solr-rank-service	634.1ms				
30	computeddataserv..._user_channel.32	solr-rank-service			8.862ms		
52	computeddataserv...ssage.64.bytekey	solr-rank-service			13.09ms		
4	computeddataserv...onical_channel.4	solr-rank-service			10.37ms		
4	computeddataserv..._canonical_user.4	solr-rank-service			8.512ms		
16	computeddataserv...ical_user_user.16	solr-rank-service			13.41ms		
52	computeddataserv...ssage.64.bytekey	solr-rank-service			16.89ms		

Blocking SolrJ



```
ServiceRequestContext ctx = RequestContext.current();  
ctx.blockingTaskExecutor().submit(() -> {  
    ...  
});
```

Blocking SolrJ



```
ServiceRequestContext ctx = RequestContext.current();  
ctx.blockingTaskExecutor().submit(() -> {  
    ...  
});
```

Blocking SolrJ

```
serverBuilder.blockingTaskExecutor(  
    new ThreadPoolExecutor(  
        size,  
        size,  
        60,  
        TimeUnit.SECONDS,  
-    new LinkedTransferQueue<>(),  
+    new SynchronousQueue<>(),  
        new DefaultThreadFactory("armeria-common-blocking-tasks", true)));
```

Blocking SolrJ

```
public class CircuitBreakingSolrClient extends SolrClient {

    public CircuitBreakingSolrClient(SolrClient solrClient, String name) {
        this.solrClient = solrClient;
        // https://line.github.io/armeria/client-circuit-breaker.html
        this.circuitBreaker = new CircuitBreakerBuilder(name)
            ...
            .build();
    }

    public NamedList<Object> request(SolrRequest solrRequest, String collection)
        throws SolrServerException, IOException {
        if (!circuitBreaker.canRequest()) {
            throw new IOException("Circuit breaker fail fast: " + circuitBreaker.name());
        }
        try {
            ...
        } catch (Exception e) {
            circuitBreaker.onFailure();
            throw e;
        }
        circuitBreaker.onSuccess();
        ...
    }
}
```

Requirements

- ✓ **Written in Java**
 - Solr, SolrJ, XGBoost4J
- ✓ **Supports multiple protocols**
 - HTTP, gRPC, Thrift-over-HTTP
- ✓ **Integrates with existing tools**
 - Consul, Prometheus, Rsyslog, Honeycomb
- ✓ **Good parallel performance**



<http://line.github.io/armeria/>

What's Next?

Next steps

- **More visibility (logging, metrics, tracing)**
- **Async Solr client**
- **Integration with Envoy Service Mesh**
- **Deployment on Kubernetes**

Thank you!
Questions?
